

# Structured Query Language (SQL) Checkpoint

This checkpoint is designed to test your understanding of writing queries in the Structured Querying Language (SQL).

Specifically, this checkpoint will cover:

- Reading an Entity Relationship Diagram
- Writing a query to return specific columns from a SQL database
- Writing a query to filter the rows from a SQL database
- Sorting observations using SQL
- Joining tables using SQL

## Data Understanding

In this repository under the file path `Northwind.sqlite` there is a SQLite database file containing information about the fictional trading company "Northwind Traders".

The tables of interest for this checkpoint will be:

**Product:** A table containing information about products sold by Northwind Traders.

**Order:** A table containing high level information about an order submitted to Northwind Traders.

**Shipper:** A table containing information about the shipping companies Northwind Traders employ to handle the shipping of their products.

## Requirements

1. Select an entire table.
2. Select all columns. Filter the rows.
3. Select a single column. Filter the rows using two conditions.
4. Sort in descending order. Return the first five rows.
5. Join two tables. Filter the rows.

## Setup

This checkpoint will test the resulting data each of your SQL queries generate. For each requirement, your query should be written as a string, and assigned to the requested variable name. The tests do not inspect the casing or formatting of your SQL query.

In the cell below we import relevant libraries.

```
# Run this cell without changes
import sqlite3
import pandas as pd
```

In the cell below we...

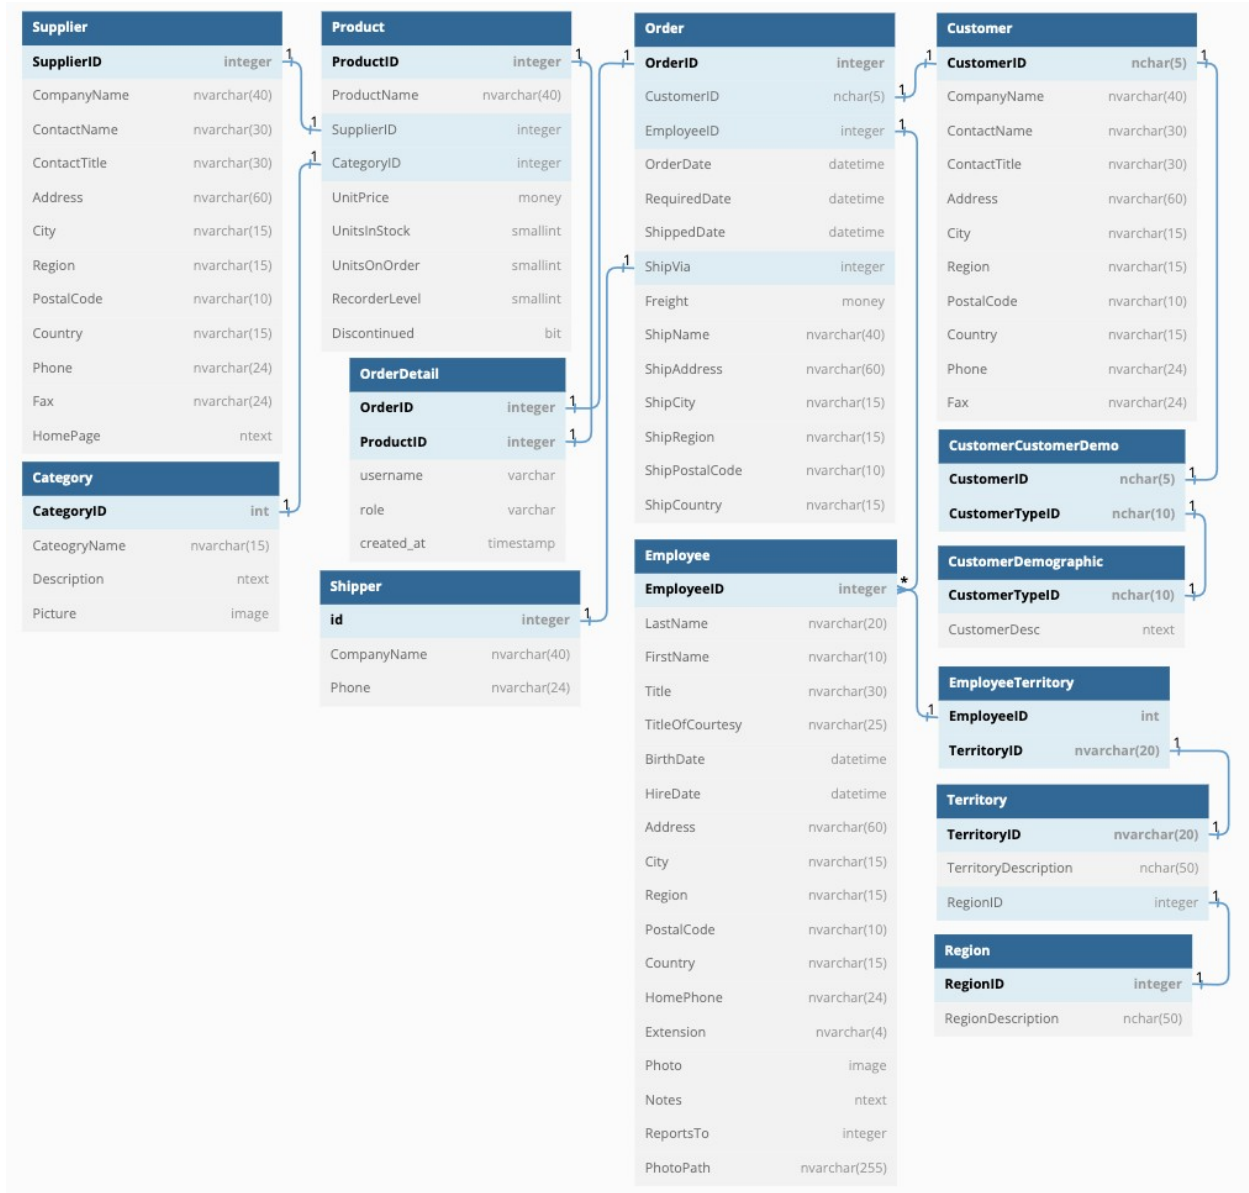
- Open up a connection to the SQLite database and store the connection in a variable called `conn`
- Initialize a SQLite cursor object with the variable name `cursor`.

```
# Run this cell without changes
northwind_path = 'Northwind.sqlite'

# Open up a connection
conn = sqlite3.connect(northwind_path)
# Initialize a cursor
cursor = conn.cursor()
```

### **Below is an Entity Relationship Diagram for the Northwind Database**

The text is quite small in the below image. Here is a [link](#) to the raw image file, where the text is slightly larger.



## Table Names

Below, we use `pd.read_sql` to output the table names in the SQLite database. When writing your queries, you should use the table names as listed below.

```
# Run this cell without changes
table_name_query = """SELECT name
                        AS 'Table Names'
                        FROM sqlite_master
                        WHERE type='table';"""

pd.read_sql(table_name_query, conn)
```

	Table Names
0	Employee
1	Category
2	Customer
3	Shipper
4	Supplier
5	Order
6	Product
7	OrderDetail
8	CustomerCustomerDemo
9	CustomerDemographic
10	Region
11	Territory
12	EmployeeTerritory

## 1. Select an entire table.

In a string variable named `first_query`, write a SQL query that returns all rows and columns from the `product` table.

When passed into `pd.read_sql` this query should return a dataframe with a head that looks like this:

		Product Name	SupplierID	CategoryID	Quantity Per Unit	Unit Price	Units In Stock	Units On Order	Reorder Level	Discontinued
0	1	Chai	1	1	10 boxes x 20 bags	18.00	39	0	10	0
1	2	Chang	1	1	24 - 12 oz bottles	19.00	17	40	25	0
2	3	Aniseed Syrup	1	2	12 - 550 ml bottles	10.00	13	70	25	0
3	4	Chef Anton's Cajun Seasoning	2	2	48 - 6 oz jars	22.00	53	0	0	0
4	5	Chef Anton's Gumbo Mix	2	2	36 boxes	21.35	0	0	0	1

```
# CodeGrade step1
# Replace None with appropriate code

first_query = """SELECT * FROM Product;"""

# Use the line below to check your query's output
pd.read_sql(first_query, conn).head()
```

	Id	ProductName	SupplierId	CategoryId	\
0	1	Chai	1	1	
1	2	Chang	1	1	
2	3	Aniseed Syrup	1	2	
3	4	Chef Anton's Cajun Seasoning	2	2	
4	5	Chef Anton's Gumbo Mix	2	2	

	QuantityPerUnit	UnitPrice	UnitsInStock	UnitsOnOrder	ReorderLevel	\
0	10 boxes x 20 bags	18.00	39	0	10	
1	24 - 12 oz bottles	19.00	17	40	25	
2	12 - 550 ml bottles	10.00	13	70	25	
3	48 - 6 oz jars	22.00	53	0	0	
4	36 boxes	21.35	0	0	0	

Discontinued
0
1
2
3
4

```

# first_query should be a string
assert type(first_query) == str

# first_query should be a SQL query
first_query_df = pd.read_sql(first_query, conn)

```

## 2. Select all columns. Filter the rows.

In a string variable named `second_query`, write a SQL query that returns all columns and rows from the `Product` table where `discontinued` has a value of `1`.

When passed into `pd.read_sql` the query's resulting dataframe should look like this:

	Id	ProductName	SupplierId	CategoryId	QuantityPerUnit	UnitPrice	UnitsInStock	UnitsOnOrder	ReorderLevel	Discontinued
0	5	Chef Anton's Gumbo Mix	2	2	36 boxes	21.35	0	0	0	1
1	9	Mishi Kobe Niku	4	6	18 - 500 g pkgs.	97.00	29	0	0	1
2	17	Alice Mutton	7	6	20 - 1 kg tins	39.00	0	0	0	1

Id	ProductName	SupplierId	CategoryId	QuantityPerUnit	UnitPrice	UnitsInStock	UnitsOnOrder	ReorderLevel	Discontinued
32	Guaraná Fantástica	10	1	12 - 355 ml cans	4.50	20	0	0	1
42	Rössle Sauerkraut	12	7	25 - 825 g cans	45.60	26	0	0	1
52	Thüringer Rostbratwurst	12	6	50 bags x 30 sausgs.	123.79	0	0	0	1
64	Singaporean Hokkien Fried Mee	20	5	32 - 1 kg pkgs.	14.00	26	0	0	1
75	Perth Pasties	24	6	48 pieces	32.80	0	0	0	1

```
# CodeGrade step2
# Replace None with appropriate code
```

```
second_query = """
SELECT * FROM Product
WHERE
    Discontinued = 1;
"""
```

```
# Use the line below to check your query's output
pd.read_sql(second_query, conn)
```

	Id	ProductName	SupplierId	CategoryId	\
0	5	Chef Anton's Gumbo Mix	2	2	
1	9	Mishi Kobe Niku	4	6	
2	17	Alice Mutton	7	6	
3	24	Guaraná Fantástica	10	1	
4	28	Rössle Sauerkraut	12	7	
5	29	Thüringer Rostbratwurst	12	6	
6	42	Singaporean Hokkien Fried Mee	20	5	
7	53	Perth Pasties	24	6	

	QuantityPerUnit	UnitPrice	UnitsInStock	UnitsOnOrder	ReorderLevel	\
0	36 boxes	21.35	0	0		
0						
1	18 - 500 g pkgs.	97.00	29	0		
0						
2	20 - 1 kg tins	39.00	0	0		
0						
3	12 - 355 ml cans	4.50	20	0		
0						
4	25 - 825 g cans	45.60	26	0		

```

0
5 50 bags x 30 sausgs.      123.79      0      0
0
6      32 - 1 kg pkgs.      14.00      26      0
0
7      48 pieces      32.80      0      0
0

    Discontinued
0          1
1          1
2          1
3          1
4          1
5          1
6          1
7          1

# second_query should be a string
assert type(second_query) == str

# second_query should be a SQL query
second_query_df = pd.read_sql(second_query, conn)

```

### 3. Select a single column. Filter the rows using two conditions.

In a string variable named `third_query`, write a SQL query that returns the name of a product if the product is not in stock and is not discontinued.

When passed into `pd.read_sql` the query's resulting dataframe should look like this:

	ProductName
0	Gorgonzola Telino

```

q = 'select Discontinued from Product;'
pd.read_sql(q, conn).Discontinued.unique()

array([0, 1], dtype=int64)

# CodeGrade step3
# Replace None with appropriate code

third_query = """
    SELECT ProductName FROM Product
    WHERE
        UnitsInStock = 0
        AND
        Discontinued = 0;

```

```

"""

# Use the line below to check your query's output
pd.read_sql(third_query, conn)

    ProductName
0  Gorgonzola Telino

# third_query should be a string
assert type(third_query) == str

# third_query should be a SQL query
third_query_df = pd.read_sql(third_query, conn)

```

## 4. Sort in descending order. Return the first five rows.

In a string variable named `fourth_query`, write a SQL query that returns the product name and unit price.

- Order by unit price in descending order.
- Use a SQL command so only the first five rows are queried.

When passed into `pd.read_sql`, this query's resulting dataframe should look like this:

	ProductName	UnitPrice
0	Côte de Blaye	263.50
1	Thüringer Rostbratwurst	123.79
2	Mishi Kobe Niku	97.00
3	Sir Rodney's Marmalade	81.00
4	Carnarvon Tigers	62.50

```

# CodeGrade step4
# Replace None with appropriate code

fourth_query = """
    SELECT
        ProductName, UnitPrice
    FROM Product
    ORDER BY
        UnitPrice DESC
    LIMIT 5;
"""

# Use the line below to check your query's output
pd.read_sql(fourth_query, conn)

    ProductName  UnitPrice
0      Côte de Blaye    263.50
1  Thüringer Rostbratwurst    123.79

```



```

2      Mishi Kobe Niku      97.00
3  Sir Rodney's Marmalade   81.00
4      Carnarvon Tigers    62.50

# fourth_query should be a string
assert type(fourth_query) == str

# fourth_query should be a SQL query
fourth_query_df = pd.read_sql(fourth_query, conn)

```

## 5. Join two tables. Filter the rows.

In a string variable named `fifth_query`, write a SQL query that returns the name and phone number for shippers who have charged more than \$1000 for shipping cost.

### Hint:

- "Freight", in the Order table, represents shipping cost.
- For this question, the word "order" is both a table name AND a SQL command. To help SQL understand that you are referencing the table `order` in the database, you will need to wrap the word `order` in quotation marks.

When passed into `pd.read_sql`, the query's resulting dataframe should look like this:

	CompanyName	Phone
0	Federal Shipping	(503) 555-9931

```

# delete
q = """
SELECT * FROM 'Order'
LIMIT 1;
"""
pd.read_sql(q, conn)

```

	Id	CustomerId	EmployeeId	OrderDate	RequiredDate	ShippedDate
ShipVia \						
0	10248	VINET	5	2012-07-04	2012-08-01	2012-07-16
3						

	Freight	ShipName	ShipAddress	ShipCity	\
0	32.38	Vins et alcools Chevalier	59 rue de l'Abbaye	Reims	

	ShipRegion	ShipPostalCode	ShipCountry
0	Western Europe	51100	France

```

# CodeGrade step5
# Replace None with appropriate code

fifth_query = """
SELECT

```

```
        Shipper.CompanyName, Shipper.Phone
FROM Shipper
JOIN 'Order'
    ON Shipper.Id = 'Order'.ShipVia
WHERE
    'Order'.Freight > 1000;
"""

# Use the line below to check your query's output
pd.read_sql(fifth_query, conn)
```

```
        CompanyName      Phone
0  Federal Shipping  (503) 555-9931

# fifth_query should be a string
assert type(fifth_query) == str

# fifth_query should be a SQL query
fifth_query_df = pd.read_sql(fifth_query, conn)
```